Daniel Finlay: dsf77

Khaled Abdeljaber: ksa68

Daniel Hollasch: djh201

ILAB: Singleton

In our implementation we used three queues, one for each level of priority, being default, high, and real-time priorities. Depending on how long a specific job ran for(interrupted or ran to completion), it went back to the beginning of the same queue, with the same priority “n”, or to the beginning of a queue with priority “n-1”, respectively. Our implementation was chosen to work in such a way that would avoid starvation of processes.

The scheduling algorithm implemented in this project is that of the sender/consumer. Where the sender thread schedules the tasks and the consumer thread executes them. So essentially, the consumer thread is like a slave, it only executes when its master(sender) thread tells it to execute a task.