Daniel Finlay: dsf77
Khaled Abdeljaber: ksa68
Daniel Hollasch: djh201

ILAB: Singleton

QUESTION 1:
The number of queues, priority levels and how much to increase and decrease priority levels is up to you. Whatever design you pick should be documented in you readme, along with your rationale.

ANSWER:
In our implementation we used three queues, one for each level of priority, being default, high, and real-time priorities. Depending on how long a specific job ran for(interrupted or ran to completion), it went back to the beginning of the same queue, with the same priority “n”, or to the beginning of a queue with priority “n-1”, respectively. Our implementation was chosen to work in such a way that would avoid starvation of processes.

QUESTION 2:
The details of your scheduler and your testing results should be written up in a readme.pdf

ANSWER:
********WAITING TO FINISH THE LAB TO WRITE UP THIS PART********
If running out of time just delete this question

QUESTION 3:
readme.pdf that explains your code's design and operation as well as your scheduling parameters.

ANSWER:
The scheduling algorithm implemented in this project is that of the sender/consumer. Where the sender thread schedules the tasks and the consumer thread executes them. So essentially, the consumer thread is like a slave, it only executes when its master(sender) thread tells it to execute a task.